

POTENTIAL-GUIDED PLANNING IN MAGIC TOWER: A DETERMINISTIC LONG-HORIZON RESOURCE-ALLOCATION BENCHMARK

Bingsong Tong (2025310812), Zhizhou Qu (2025310268), Rui Chen (2025211105)

Deep Reinforcement Learning Course Project 2026

ABSTRACT

Magic Tower is a deterministic single-agent role-playing puzzle in which the agent must manage hit points (HP), attack (ATK), defense (DEF), keys, and money across many irreversible interactions. We study the first-ten-floor task, which is fully observable and rule-defined, yet naive reinforcement learning rarely receives a useful terminal signal. The deterministic combat formula creates sharp resource thresholds: a locally legal interaction may preserve immediate feasibility while eliminating a much later solution path. We organize the study around two complementary approaches. First, agent-assisted reference-route construction produces high-quality trajectories for auditing simulator semantics and revealing feasible resource schedules. Second, an AlphaGo Zero-style learning-and-planning pipeline combines a graph policy/value model, single-agent PUCT-MCTS, and potential-based reward shaping learned from reference trajectories. The reference-aided policy/value variant reaches the terminal condition from the fixed initial state with final HP 315, while the no-demonstration variant remains incomplete, suggesting that local search priors must be combined with archive-based exploration and continuation-aware stage values. A separate multi-agent reference planner, based on specialized candidate scorers and beam search over legal simulator transitions, obtains a validated trajectory with final HP 436. This result is reported as a reference-construction baseline rather than as evidence for no-demonstration exploration.

1 INTRODUCTION

Magic Tower is a grid-based role-playing puzzle in which the agent ascends a multi-floor tower by navigating corridors, collecting resources, opening doors, triggering events, and fighting enemies. Each floor is a discrete map whose cells contain walls, stairs, consumable items, key-locked doors, enemies, or non-player events. The agent’s state is summarized by its location and a small set of persistent resources: HP, ATK, DEF, money, and colored keys. A single interaction may consume a map object, change the reachable region of the tower, or permanently modify the agent’s combat statistics. Progress is therefore not determined by local navigation alone, but by the order in which the agent converts limited resources into future accessibility.

The game is fully observable and deterministic, but this does not make it trivial for reinforcement learning. Combat damage is computed by a fixed formula from the agent and enemy statistics, so one additional point of ATK or DEF can abruptly change the cost of multiple future encounters. Doors and consumable items introduce irreversible resource allocation: opening a locally useful door may spend the key needed to access a later damage-reducing item, while taking an HP potion too early may remove the buffer needed for a more efficient future route. As a result, Magic Tower is naturally modeled as a deterministic long-horizon planning problem over a dynamically changing resource-interaction graph.

We study a first-ten-floor instance in which success requires reaching the terminal encounter with positive HP under deterministic validation. The task contains the core dependencies of the full game: early equipment acquisition, key circulation, threshold-sensitive stat growth, and a terminal combat check. This paper develops two complementary solution paths. First, agent-assisted reference-




Figure 1: Representative floors from the first-ten-floor Magic Tower instance. Each panel is a tile-rendered floor map: traversable corridors connect stairs, resource pickups, doors, enemies, and event triggers. The figure illustrates why the task is a multi-floor resource-interaction problem rather than a purely local navigation problem.

route construction is used to obtain high-quality trajectories and to validate simulator semantics. A multi-agent beam-search planner constructs a strong reference trajectory with final HP 436, while an earlier human-audited reference route with final HP 403 is retained for potential learning and simulator audits. Second, we build an AlphaGo Zero-style learning-and-planning system that combines graph policy/value learning, single-agent PUCT-MCTS, and potential-based reward shaping. The reference-aided policy/value variant also solves the first-ten-floor instance from the fixed initial state, reaching final HP 315, whereas the no-demonstration variant still reaches only intermediate milestones. This gap motivates archive-based exploration and continuation-aware stage values.

2 RELATED WORK

Search-guided reinforcement learning. AlphaGo Zero and AlphaZero demonstrated that policy/value networks can be improved through Monte Carlo tree search without relying on handcrafted evaluation functions (Silver et al., 2017; 2018). These systems train a network $f_\theta(s) = (p_\theta(\cdot | s), v_\theta(s))$, use MCTS to convert the raw policy prior into an improved visit-count policy, and then regress the network back toward the search policy and realized return. MCTS itself can be viewed as a repeated selection, expansion, evaluation, and backup procedure whose action rule balances exploitation and uncertainty (Kocsis & Szepesvári, 2006; Browne et al., 2012). In AlphaZero-style PUCT, this balance is further biased by a learned prior over legal actions. These systems, however, are designed for two-player board games with frequent terminal supervision and alternating max/min structure. Magic Tower is a single-agent deterministic MDP: every decision layer is a maximization layer, MCTS backups accumulate discounted task return without a sign flip, and random interaction rarely reaches a successful terminal state. MuZero further showed that planning can be performed with a learned dynamics model (Schrittwieser et al., 2020); in this project the transition model is known, so the central question is not model learning but how to combine symbolic simulation, policy/value priors, and long-horizon search.

Hard exploration and archive methods. Magic Tower belongs to the hard-exploration regime because most early action sequences either exhaust resources or enter irreversible dead states before any useful terminal reward is observed. Go-Explore addresses this failure mode by remembering promising intermediate states, returning to them, and exploring outward from those stepping stones (Ecoffet et al., 2019). This idea is well aligned with Magic Tower: a state after obtaining a key item, preserving a key buffer, or satisfying a combat threshold is often more valuable than a locally high-scoring route prefix. The archive-MCTS component follows this perspective by treating stage states as reusable search anchors rather than restarting all exploration from the initial state. Because the simulator is deterministic, returning to an archived state can be implemented by storing the state object or by replaying its route prefix; the later robustification idea in Go-Explore is analogous to training a policy/value model from successful archive trajectories, but is kept separate from pure exploration in this report.

Reward shaping and demonstrations. Sparse terminal rewards make direct reinforcement learning inefficient in Magic Tower. Potential-based reward shaping provides a principled way to add

dense feedback while preserving the optimal policy under suitable conditions (Ng et al., 1999). Inverse reinforcement learning and maximum-entropy trajectory modeling offer broader tools for inferring reward structure from demonstrations (Ziebart et al., 2008). Our use of successful trajectories is narrower: reference routes define an ordered sequence of states from which a stage-conditioned potential is estimated. This potential supplies dense search rewards, while final evaluation remains tied to deterministic task success. The important distinction is that the potential term $\gamma\Phi_\phi(s', g') - \Phi_\phi(s, g)$ changes intermediate learning signals through a telescoping difference, whereas the terminal task criterion remains unchanged.

Puzzle, roguelike, and reasoning benchmarks. Sokoban-style environments emphasize irreversible spatial constraints and generalization across levels (Justesen et al., 2018), while NetHack and MiniHack emphasize symbolic observations, complex entity interactions, and long-horizon decision making (Küttler et al., 2020). Magic Tower sits between these families: it is fully observable and deterministic like many puzzle domains, yet its difficulty is dominated by numerical resource thresholds and multi-floor dependency structure. ARC-AGI is also relevant because it frames intelligence as abstract rule inference and test-time adaptation (Chollet, 2019; Chollet et al., 2024; 2025). The interactive ARC-AGI-3 setting extends this line toward agents that infer goals and dynamics through action (ARC Prize Foundation, 2026; Rodionov, 2026). Magic Tower offers a complementary benchmark: it has explicit transition dynamics, textual and video-renderable trajectories, and a large community corpus of demonstrations and play episodes.

3 PROBLEM MODELING AND SEARCH INTERFACE

3.1 DETERMINISTIC MDP

We model the task as a finite-horizon deterministic single-agent MDP

$$\mathcal{M} = (\mathcal{S}, \mathcal{A}_{\text{valid}}, f, r_{\text{task}}, \gamma, \rho_0, \mathcal{G}, \mathcal{F}, H).$$

All experiments use a fixed start state, so ρ_0 is a point mass. A state $s \in \mathcal{S}$ contains the complete map state of Floors 1–10, the agent floor and coordinate, hero resources, key counts, consumed objects, triggered events, and Boolean event variables. We write $\text{HP}(s)$, $\text{ATK}(s)$, $\text{DEF}(s)$, and $\text{Money}(s)$ for the four numeric hero resources. The valid action set is the set of executable abstract actions:

$$a_t \in \mathcal{A}_{\text{valid}}(s_t).$$

Each action chooses an interaction node, not an individual arrow key. The simulator expands it into a shortest path to a reachable item, stair, event, merchant, openable door, or killable enemy. The transition is deterministic:

$$s_{t+1} = f(s_t, a_t).$$

Let $b_{\text{term}}(s) \in \{0, 1\}$ denote the terminal event variable set by the simulator after the Floor-10 terminal adversary is defeated. Terminal success is the state set

$$\mathcal{G} = \{s \in \mathcal{S} : b_{\text{term}}(s) = 1 \wedge \text{HP}(s) > 0\}.$$

Failures include death, invalid transitions, timeouts, and dead states with no useful executable actions; by construction $\mathcal{G} \cap \mathcal{F} = \emptyset$. The sparse task reward is

$$r_{\text{task}}(s, a, s') = R_G \mathbf{1}[s' \in \mathcal{G}] - R_F \mathbf{1}[s' \in \mathcal{F}] - \kappa,$$

where κ is a small abstract-step cost. For terminal states we set $V^*(s) = 0$. For non-terminal states, the optimal value satisfies the single-player Bellman equation

$$V^*(s) = \max_{a \in \mathcal{A}_{\text{valid}}(s)} [r_{\text{task}}(s, a, f(s, a)) + \gamma V^*(f(s, a))].$$

This yields an AlphaGo Zero-style search interface, but the target is a deterministic single-agent return rather than an adversarial game outcome. A PUCT-MCTS rollout backs up

$$G_k = r_k + \gamma G_{k+1} = \sum_{i=k}^{L-1} \gamma^{i-k} r_i + \gamma^{L-k} \hat{V}(s_L).$$

The search policy at the root is obtained from visit counts,

$$\pi_{\text{MCTS}}(a \mid s) = \frac{N(s, a)^{1/\tau}}{\sum_{b \in \mathcal{A}_{\text{valid}}(s)} N(s, b)^{1/\tau}},$$

and an AlphaZero-style single-agent learner would train on triples $(s_t, \pi_{\text{MCTS}}(\cdot \mid s_t), G_t)$, replacing the two-player win/loss label with cumulative task return.

Combat is the main discontinuity. For an ordinary enemy e , let (H_e, A_e, D_e) be enemy HP, ATK, and DEF, and let $(A_h(s), D_h(s))$ be hero ATK and DEF. If $A_h(s) \leq D_e$, the enemy is not killable. Otherwise, the number of hero attacks is

$$T(s, e) = \left\lceil \frac{H_e}{A_h(s) - D_e} \right\rceil,$$

and the ordinary combat loss implemented by the simulator is

$$C(s, e) = (T(s, e) - 1) \max(0, A_e - D_h(s)).$$

The factor $T - 1$ reflects that the enemy does not counterattack after the final hero strike. The simulator adds the configured first-strike, fixed-damage, counterattack, and armor-breaking terms for enemies that carry special effects. Thus a single red or blue gem can reduce future damage by hundreds of HP, while the same gem may have no effect away from a threshold. This makes dense reward design unusually sensitive.

3.2 SIMULATOR AND MACRO-ACTION INTERFACE

The simulator is implemented in Python and uses the extracted HTML5 data as the truth source. It implements deterministic combat, resource collection, doors, stairs, merchants, and event transitions required by the first-ten-floor instance. The visualizer is kept as an independent validation interface; search and learning interact only with the simulator state.

Primitive arrow-key control is unnecessarily long-horizon for planning. Instead, the agent chooses an interaction node. For a state s on floor f , let $G_f = (V_f, E_f)$ be the precomputed static adjacency graph of walkable grid cells, and let $B_f(s) \subset V_f$ be the dynamic blocker set containing currently closed doors, undefeated enemies, blocking events, and non-pass-through NPC cells. Reachability is computed on the induced graph

$$G_f(s) = G_f[V_f \setminus B_f(s)].$$

A single BFS from the hero position gives the reachable set $R_f(s)$, distances $\text{dist}_s(v)$, and a parent pointer

$$\text{par}_s(v) = (u, \delta), \quad u \in R_f(s), \delta \in \{\uparrow, \downarrow, \leftarrow, \rightarrow\}.$$

The macro-action set is then generated in two passes. First, every reachable interaction node is executable immediately. Second, each blocker adjacent to $R_f(s)$ becomes a candidate door, enemy, or event action, and a local legality predicate checks keys, battle survivability, money, and event flags:

$$\mathcal{A}_{\text{valid}}(s) = \{a_i : i \in I_f, i \in R_f(s)\} \cup \{a_j : j \in \partial R_f(s), L(s, j) = 1\}.$$

This abstraction keeps ordinary movement inside the transition function while preserving all resource-changing decisions and floor-transfer choices as abstract actions.

The expensive part is not BFS itself, since each floor has only a small fixed grid; the cost comes from repeating it inside thousands of tree expansions and allocating full paths for targets that are never selected. We therefore use a fast reachability layer: static adjacency lists and interaction-node tables are precomputed once, $B_f(s)$ is encoded as a compact bitset, BFS results are cached by $(f, \text{pos}(s), B_f(s))$, and paths are reconstructed from par_s only for the chosen action. Scalar variables such as HP, ATK, DEF, money, and keys are excluded from the BFS cache key; they affect only the local legality predicate $L(s, j)$, not geometric passability. This separation reduces CPU time in MCTS by turning most action enumeration calls into bitset lookup, cached reachability, and cheap local legality checks.

3.3 EVALUATION SETTING AND REFERENCE TRAJECTORIES

For validation and later comparison, we construct high-quality reference trajectories with external planning assistance and strict simulator replay. The purpose is to expose feasible resource schedules, validate the simulator, and establish reference behavior for subsequent reward and planning experiments. These trajectories verify deterministic transition semantics, event logic, no-shop/no-fly constraints, and visualization consistency. The fixed start state is Floor 2 at coordinate (3, 7), with $HP = 400$, $ATK = 10$, $DEF = 10$, $Money = 4$, and zero colored keys. The reference trajectory reaches the Floor-10 terminal condition with

$$HP_T = 403, \quad ATK_T = 27, \quad DEF_T = 27, \quad Money_T = 262.$$

It uses no Floor-4 shop and no flying item. The trajectory is retained as a validation artifact, not as evidence that the pure RL system has solved the task.

4 METHODOLOGY

4.1 REFERENCE-DERIVED POTENTIAL LEARNING

The reference trajectory is too narrow to define the full optimal policy. We therefore use it as an ordered sequence of states for estimating a potential function $\Phi_\phi(s, g)$, following the policy-invariance motivation of potential-based reward shaping (Ng et al., 1999). For a successful trajectory $\tau = (s_0, a_0, \dots, s_T)$, we create progress and ranking labels:

$$y_t = t/T, \quad s_j \succ s_i \quad \text{if } j > i.$$

The potential loss combines regression, pairwise ranking, and dead-end separation:

$$\begin{aligned} \mathcal{L}_\Phi(\phi) = & \sum_t (\Phi_\phi(s_t, g_t) - y_t)^2 \\ & - \alpha \sum_{j>i} \log \sigma(\Phi_\phi(s_j, g_j) - \Phi_\phi(s_i, g_i)) \\ & - \beta \sum_{(s^+, s^-)} \log \sigma(\Phi_\phi(s^+, g^+) - \Phi_\phi(s^-, g^-)). \end{aligned}$$

Let $g(s)$ denote the deterministic stage label associated with state s . The search reward is then potential-based. With $g = g(s)$ and $g' = g(s')$ denoting the stage labels before and after the transition,

$$r_{\text{search}}(s, g, a, s', g') = r_{\text{task}}(s, a, s') + \lambda (\gamma \Phi_\phi(s', g') - \Phi_\phi(s, g)).$$

Along a finite trajectory, the discounted sum of the shaping term telescopes:

$$\sum_{t=0}^{T-1} \gamma^t (\gamma \Phi_\phi(s_{t+1}, g_{t+1}) - \Phi_\phi(s_t, g_t)) = -\Phi_\phi(s_0, g_0) + \gamma^T \Phi_\phi(s_T, g_T).$$

Here $g_t = g(s_t)$, so the stage-conditioned potential is equivalently a potential over an augmented Markov state. Strict policy invariance follows in the standard absorbing-state formulation, or when terminal-state potentials are fixed so that the final term does not distinguish successful terminal states. In this project the shaped reward is therefore treated as a search and learning signal rather than as the evaluation objective. Thus the potential can bias local learning and search toward states that look closer to success, while the original terminal objective remains the quantity used for evaluation. This keeps task evaluation separate from shaping. Final success is always measured by deterministic validation and the terminal predicate b_{term} , not by shaped reward.

We use a linear potential over audited semantic factors rather than a free-form neural reward model. Let $c_k(s; g)$ denote the k -th stage-conditioned component for stage g , including agent resources, key pressure, reachable resources, blocked resource pressure, monster damage, marginal damage decrease under $ATK+1$ or $DEF+1$, equipment progress, Floor-10 resource progress, and terminal-feasibility margin. The learned potential is

$$\Phi_\phi(s, g) = \sum_k w_{g,k} c_k(s; g) + \sum_k w_k^{\text{global}} c_k(s; g).$$

For later notation, let $\bar{w}_{g,k} = w_{g,k} + w_k^{\text{global}}$ denote the effective component weight.

The reference trajectory is used to estimate only the weights $w_{g,k}$, not the feature definitions. Each component c_k is computed from the simulator and resource graph. We group the components as

$$c(s; g) = [c^{\text{hero}}, c^{\text{key}}, c^{\text{reach}}, c^{\text{combat}}, c^{\text{threshold}}, c^{\text{stage}}, c^{\text{terminal}}].$$

The reference trajectory provides stage-specific constraints on the relative weights of these groups at different stages. The hero block contains HP, ATK, DEF, money, and key counts. The key and reachability blocks measure current key pressure, executable resources, and resources blocked by doors or enemies. The combat block contains deterministic damage to killable enemies and the terminal adversary. The threshold block contains marginal damage drops under one additional ATK or DEF:

$$\Delta_a(s, e) = C(s, e) - C(s^{a+1}, e), \quad \Delta_d(s, e) = C(s, e) - C(s^{d+1}, e),$$

where s^{a+1} and s^{d+1} denote counterfactual states with one extra ATK or DEF. The stage block records milestone progress such as equipment, stat resources, key items, and terminal-gate readiness, while the terminal block contains the final HP margin and the terminal event variable.

Training examples compare one-step potential changes among legal abstract actions. Define

$$\Delta_\phi(s, a; g) = \gamma \Phi_\phi(f(s, a), g(f(s, a))) - \Phi_\phi(s, g),$$

where the next-stage label is evaluated after applying the transition. At trajectory step t , the simulator enumerates $\mathcal{A}_{\text{valid}}(s_t)$. The reference transition $(s_t, a_t^E, f(s_t, a_t^E))$ provides a positive ranking candidate for the one-step potential difference:

$$\Delta_k^E = \gamma c_k(f(s_t, a_t^E); g(f(s_t, a_t^E))) - c_k(s_t; g_t),$$

and each valid alternative a^- provides

$$\Delta_k^- = \gamma c_k(f(s_t, a^-); g(f(s_t, a^-))) - c_k(s_t; g_t).$$

The ranker optimizes

$$\mathcal{L}_{\text{rank}} = \sum_{t, a^-} \text{softplus} \left(- \sum_k \bar{w}_{g_t, k} (\Delta_k^E - \Delta_k^-) \right) + \eta \|w\|_2^2.$$

Here $\bar{w}_{g_t, k}$ includes both the stage-specific and global contributions. Equivalently, the learned potential is trained so that

$$\Delta_\phi(s_t, a_t^E; g_t) > \Delta_\phi(s_t, a; g_t), \quad a \in \mathcal{A}_{\text{valid}}(s_t) \setminus \{a_t^E\},$$

whenever the reference action advances the successful reference route more than the legal alternative. Dead-end search states add negative pairs by requiring states on successful prefixes to have larger potential than states that have no useful continuation. Thus the reference data affect the learned shaping term through stage-specific factor weights, progress labels, and pairwise transition preferences. They do not alter the sparse terminal reward, the transition function, the action mask, or the PUCT policy prior. The output is therefore a stage-conditioned potential, not a policy prior. The downstream AlphaZero-style and archive-MCTS experiments use the reference trajectory only through


$$r_\Phi(s, g, a, s', g') = \gamma \Phi_\phi(s', g') - \Phi_\phi(s, g),$$

which enters single-agent MCTS backups as an edge reward and value-return target. The potential-only experiment uses no reference prefix or reference action prior during search.

4.2 MULTI-AGENT REFERENCE PLANNER

The reference-construction line uses the same simulator and legal macro-action interface, but decomposes candidate evaluation into specialized scoring modules. At state s_t , the simulator enumerates the legal candidate set $\mathcal{A}_{\text{valid}}(s_t)$ and computes the deterministic successor $s^a = f(s_t, a)$ for each candidate. All modules observe the same tuple

$$x(s_t, a) = (s_t, a, s^a, \mathcal{A}_{\text{valid}}(s_t), m_t, \mathcal{V}_t),$$



All decisions are made over simulator-verified legal candidates; reference alignment is used only in this planner.

Figure 2: Multi-agent reference planner. The simulator enumerates legal macro-actions and candidate after-states; specialized scoring modules evaluate the same candidate table; a beam-search arbiter combines scores, consults recent and visited-state memories, and writes a replayable route.

where m_t is a recent trace memory and $\mathcal{V}_t$ is a visited-state memory. The modules score each candidate from complementary perspectives: stage progress, key and item economy, combat thresholds, terminal-objective readiness, and optional reference-route alignment. If $q_j(s, a)$ denotes the score from module j , the arbiter uses

$$S(s, a) = \sum_j \omega_j q_j(s, a) + \omega_{\text{ref}} q_{\text{ref}}(s, a) - \lambda_{\text{loop}} \mathbf{1}[\kappa(s^a) \in \mathcal{V}_t],$$

where $\kappa(s)$ is a compact state key used for loop suppression. The reference-alignment term is used only in this reference planner; it is not part of the no-demonstration archive or pure PUCT experiments.

The arbiter performs beam search over simulator-verified transitions. Given a beam $\mathcal{B}_t = \{(s_i, \tau_i, R_i)\}_{i=1}^K$, the next beam is

$$\mathcal{B}_{t+1} = \text{TopK}_{(s_i, a)} \{(f(s_i, a), \tau_i \circ a, R_i + S(s_i, a)) : a \in \mathcal{A}_{\text{valid}}(s_i)\}.$$

The planner therefore never edits the game state directly: every candidate route is a sequence of legal simulator transitions. Recent trace memory records the latest stages, actions, and resource summaries, while visited-state memory penalizes repeated floor-transition loops and semantically equivalent revisits. This design provides a high-quality reference-route construction mechanism and diagnostic traces for later potential learning, without changing the terminal evaluation criterion.

4.3 ARCHIVE-GUIDED PUCT-MCTS

The no-demonstration exploration line combines a global archive with local MCTS. A cell abstraction maps states to coarse stepping stones:

$$\chi(s) = (\text{stage}, \text{floor}, a_b, d_b, k_b, h_b, \text{F10_progress}, \text{terminal_margin}_b).$$

The archive stores top- K non-identical candidates per cell:

$$\mathcal{B}[\chi(s)] = \{(s_i, \tau_i, \text{score}_i)\}_{i=1}^K.$$

At each iteration, the algorithm first returns to a selected cell and then explores from it. The cell is sampled according to novelty, progress, and visit penalty. From the archived state s_c , PUCT-MCTS selects a local action:

$$a = \arg \max_{a \in \mathcal{A}_{\text{valid}}(s)} \left[Q(s, a) + c_{\text{puct}} P(s, a) \sqrt{\frac{\sum_{b \in \mathcal{A}_{\text{valid}}(s)} N(s, b)}{1 + N(s, a)}} \right].$$

For the pure line P is uniform or learned from self-generated search traces, not from reference routes. In the potential-shaping line, the reference trajectory influences only the edge reward and leaf value through Φ_ϕ . The archive score is lexicographic in success, stage progress, HP, keys, and path length, so a high shaping score cannot replace task success.




Figure 3: System architecture. The simulator provides deterministic abstract transitions and resource-graph states; archive search and PUCT operate on this symbolic interface, while visualization is used only for independent trajectory inspection.




Figure 4: Summary of the main empirical signals. Left: the learned potential ranks the reference transition within the top three legal alternatives in most validated states. Middle: graph policy/value training reaches its best checkpoint before overfitting becomes visible. Right: deterministic replay separates reference validation, reference-aided closed-loop search, and the current no-demonstration archive search.

5 EMPIRICAL RESULTS

Evaluation criterion. All reported successes are verified by exact deterministic replay. A trajectory $\tau = (s_0, a_0, \dots, s_T)$ is counted as successful only if it reaches the terminal event and retains positive hit points:

$$\text{Eval}(\tau) = \mathbf{1}[b_{\text{term}}(s_T) = 1 \wedge \text{HP}(s_T) > 0].$$

Potential scores, shaped returns, supervised losses, and search scores are used only as diagnostic quantities. This separation is important because a shaped reward can improve local search without certifying that the task objective has been reached. We therefore distinguish reference validation, potential learning, reference-aided policy/value search, perturbation diagnostics, and no-demonstration archive search.

Potential learning. The potential model is trained from ordered states rather than from direct action imitation. For each validated transition we compare the reference successor s^+ with alternative legal successors s^- , and optimize a pairwise ranking objective

$$\mathcal{L}_\Phi(\phi) = - \sum_{(s^+, s^-)} \log \sigma(\Phi_\phi(s^+, g^+) - \Phi_\phi(s^-, g^-)).$$

The resulting dataset contains 317 validated states and 1578 pairwise comparisons. The learned potential ranks the reference transition first in 52.1% of states and within the top three in 85.8%, with mean rank 2.00. These numbers measure shaping quality rather than policy accuracy: Φ_ϕ only biases search through $\gamma\Phi_\phi(s', g') - \Phi_\phi(s, g)$, while the sparse terminal criterion remains unchanged.




Figure 5: Route structure of the multi-agent reference planner. Left: checkpoint timing along the macro-action sequence. Right: checkpoint-level HP, ATK, and DEF values, showing the trade-off between immediate hit-point loss and later combat-threshold improvement.

Graph policy/value model. The reference-aided comparison uses 9 unique successful trajectories after exact action-sequence deduplication, yielding 2576 graph-state supervision samples. The model is a six-layer Transformer with $d_{\text{model}} = 512$, eight attention heads, dropout 0.03, batch size 512, and AdamW learning rate $8 \cdot 10^{-5}$. It maps a graph state and stage token to a masked action prior and a scalar value estimate:

$$f_{\theta}(G_s, g) = (p_{\theta}(\cdot \mid s, g), v_{\theta}(s, g)),$$

trained with

$$\mathcal{L}(\theta) = -\pi^{\top} \log p_{\theta}(\cdot \mid s, g) + 0.35 (G - v_{\theta}(s, g))^2 + 10^{-4} \|\theta\|_2^2.$$

The best checkpoint occurs at step 9100, with total loss 0.0064 and policy loss 0.0063. The subsequent increase in policy loss is consistent with the limited diversity of the successful-trajectory set. We therefore use the model as a search prior inside PUCT rather than treating it as an autonomous policy.

Multi-agent reference planning. The multi-agent reference planner produces a legal trajectory of 291 macro-actions. Under strict replay, the terminal event is reached at the final Floor-10 encounter with

$$\text{HP}_T = 436, \quad \text{ATK}_T = 27, \quad \text{DEF}_T = 27.$$

The route passes the same deterministic replay and constraint checks as the other reported routes. Its milestone sequence is useful because it exposes the long-horizon structure that local greedy search repeatedly misses: the route obtains the sword at step 19, obtains the shield at step 75, collects late Floor-10 stat resources before the terminal stage, obtains the red key at step 248, opens the terminal red door at step 281, triggers the Floor-10 mechanism at step 282, and reaches terminal success at step 290. The resource curves show that the route does not monotonically maximize HP; instead, it accepts controlled HP losses to improve ATK and DEF, then rebuilds the HP margin before the terminal sequence.

Closed-loop validation. The policy/value model is useful only when coupled with search over the complete legal abstract-action set. Stage-restricted action filters can remove recovery actions after small deviations and thereby obscure the value of the learned prior. With the full legal action set, PUCT solves the first-ten-floor task from the fixed initial state, ending with $\text{HP}_T = 315$, $\text{ATK}_T = 27$, and $\text{DEF}_T = 27$. Additional terminal-HP sweeps with the same architecture produced successful trajectories, but did not improve upon this checkpoint. The result supports a narrow conclusion: the learned prior is an effective guide on the original instance, but the experiment does not yet establish robust policy learning.

Generalization and exploration diagnostics. Two diagnostic tests expose the current limitations. First, local item-swap perturbations show that the model is sensitive to object relocation: when a late stat resource is exchanged with a nearby key, closed-loop PUCT follows the training-distribution route structure and does not reliably recover the moved stat resource. Second, PUCT can solve several suffix subproblems when initialized from meaningful intermediate states, but independently

Condition	Evidence source	Evaluation	Outcome
Reference validation	Human-audited route	Strict replay	Terminal event reached with $HP_T = 403$; used to validate simulator semantics and reward factors.
Multi-agent reference planner	Specialized scorers and beam arbiter	Strict replay	Terminal event reached in 291 macro-actions with $HP_T = 436$, $ATK_T = 27$, and $DEF_T = 27$.
Potential shaping	Ordered states	Pairwise ranking	Rank@1 is 52.1%, Rank@3 is 85.8%, and mean rank is 2.00.
Policy/value search	Successful routes	PUCT	Closed-loop search reaches the terminal event from the fixed initial state with $HP_T = 315$.
Layout perturbation	Modified resource layout	PUCT	The model does not reliably recover a relocated late-stage stat resource.
No-demonstration archive	Simulator-generated traces	Archive-MCTS	Several milestones are reached, but terminal success has not yet been obtained.

Table 1: Primary empirical outcomes. Success is reported only under strict deterministic replay; all other quantities are diagnostic signals used to interpret shaping, search, and generalization behavior.

Diagnostic	Result	Interpretation
Full-action PUCT	terminal success with $HP_T = 315$	Search benefits from learned priors when the legal mask preserves recovery actions.
Suffix planning	local suffixes solved	Intermediate-state reuse is necessary for composing long routes.
Stage composition	terminal event not reached	Stage values must encode future feasibility, not only immediate resource completion.
Resource relocation	late stat resource missed	The current model relies too strongly on layout-specific route structure.
No-demonstration archive	incomplete terminal progress	Exploration remains the central bottleneck without reference-derived signals.

Table 2: Diagnostic experiments. The table separates what is already solved by search-guided priors from the remaining weaknesses in stage composition, object-relocation generalization, and no-demonstration exploration.

generated stage states do not compose reliably. The dominant failure is continuation awareness: a locally successful stage may consume the key or HP buffer required by a later stage.

Finally, profiling shows that planning is simulator-bound rather than neural-inference-bound. The main costs are action enumeration, reachability computation, state copying, and local legality checks. This observation is important for interpreting model scale: increasing the policy/value model size is unlikely to improve search until the symbolic expansion layer uses compact state encodings, transposition tables, reachability caches, lazy path recovery, and batched leaf evaluation.

6 FUTURE WORK: MAGIC TOWER AS AN LLM BENCHMARK

Magic Tower can be turned into a benchmark suite beyond this first-ten-floor RL project. The environment can expose three synchronized interfaces: symbolic JSON state, natural-language state summaries, and video frames from the visualizer. An LLM agent can then be evaluated on action selection, trajectory-level planning, rule discovery, and post-hoc explanation. This is analogous to ARC-AGI’s emphasis on test-time reasoning, but Magic Tower adds an explicit transition function and long interactive trajectories.

A major advantage of this domain is data scale. The Magic Tower community has produced thousands of user-created tower instances and a large number of demonstrations, walkthroughs, and

complete play episodes. This corpus can support multiple benchmark tracks: supervised trajectory modeling from successful play, offline RL from mixed-quality episodes, reward learning from ordered demonstrations, and held-out generalization across map layouts and tower designs. A robust benchmark must separate seen towers, unseen towers, rule preserving layout perturbations, and observation modalities such as symbolic state, text-only summaries, and video-only interaction. This would make Magic Tower a useful complement to ARC-style reasoning benchmarks: it preserves symbolic rule structure while requiring long-horizon interactive resource planning.

7 LIMITATIONS

This study is limited to a first-ten-floor research instance rather than the complete 50-floor game. The macro-action interface is appropriate for resource planning, but it abstracts away primitive movement and therefore does not test end-to-end pixel control. The multi-agent reference planner uses explicit scoring modules and a reference-alignment term, so its result should be viewed as a strong planning baseline and route-construction tool rather than as an autonomous no-demonstration solver. The reference-aided model is trained on a small number of successful trajectories; we therefore treat replay performance as interpolation rather than broad generalization. The perturbation experiments show that the current policy/value model is still sensitive to item relocation, and the archive experiments show that no-demonstration search has not yet discovered a complete trajectory. The main remaining technical limitations are continuation-aware value estimation, archive diversity, simulator expansion throughput, and robustness to rule-preserving layout perturbations. Finally, reference trajectories are useful for simulator validation and shaping, but must remain methodologically separate from no-demonstration reinforcement-learning claims.

8 CONCLUSION

Magic Tower is best treated as a deterministic long-horizon planning benchmark with sparse terminal reward and discontinuous resource thresholds. Pure AlphaZero-style root search is insufficient because useful terminal states are rare and early local choices can invalidate later resource schedules. The experiments support four conclusions. First, reference trajectories are useful for validating simulator semantics and calibrating reward factors, but they are methodologically distinct from autonomous exploration. Second, decomposing candidate evaluation into stage, resource, combat, and terminal-objective modules yields a strong reference-construction baseline, reaching terminal success with $HP_T = 436$ under strict replay. Third, reference-derived potential shaping can provide dense stage-progress signals while preserving sparse terminal evaluation. Fourth, reference-aided policy/value learning can solve local suffix subproblems and one unperturbed closed-loop instance, but robust stage composition remains unresolved. The central empirical lesson is that successful Magic Tower agents must reason about future resource feasibility rather than merely optimize local stage completion.

REFERENCES

- ARC Prize Foundation. ARC-AGI-3: A new challenge for frontier agentic intelligence. *arXiv preprint arXiv:2603.24621*, 2026.
- Cameron B. Browne, Edward Powley, Daniel Whitehouse, Simon M. Lucas, Peter I. Cowling, Philipp Rohlfshagen, Stephen Tavener, Diego Perez, Spyridon Samothrakis, and Simon Colton. A survey of monte carlo tree search methods. *IEEE Transactions on Computational Intelligence and AI in Games*, 4(1):1–43, 2012.
- François Chollet. On the measure of intelligence. *arXiv preprint arXiv:1911.01547*, 2019.
- Francois Chollet, Mike Knoop, Gregory Kamradt, and Bryan Landers. ARC Prize 2024: Technical report. *arXiv preprint arXiv:2412.04604*, 2024.
- Francois Chollet, Mike Knoop, Gregory Kamradt, Bryan Landers, and Henry Pinkard. ARC-AGI-2: A new challenge for frontier ai reasoning systems. *arXiv preprint arXiv:2505.11831*, 2025.

- Adrien Ecoffet, Joost Huizinga, Joel Lehman, Kenneth O. Stanley, and Jeff Clune. Go-explore: a new approach for hard-exploration problems. *arXiv preprint arXiv:1901.10995*, 2019.
- Niels Justesen, Philip Bontrager, Julian Togelius, and Sebastian Risi. Illuminating generalization in deep reinforcement learning through procedural level generation. *NeurIPS Workshop on Deep Reinforcement Learning*, 2018.
- Levente Kocsis and Csaba Szepesvári. Bandit based monte-carlo planning. In *European Conference on Machine Learning*, pp. 282–293, 2006.
- Heinrich Küttler, Nantas Nardelli, Alexander Miller, Roberta Raileanu, Marco Selvatici, Edward Grefenstette, and Tim Rocktäschel. The nethack learning environment. *NeurIPS*, 2020.
- Andrew Y. Ng, Daishi Harada, and Stuart Russell. Policy invariance under reward transformations: Theory and application to reward shaping. In *Proceedings of the Sixteenth International Conference on Machine Learning*, 1999.
- Sergey Rodionov. Executable world models for ARC-AGI-3. *arXiv preprint arXiv:2605.05138*, 2026.
- Julian Schrittwieser, Ioannis Antonoglou, Thomas Hubert, Karen Simonyan, Laurent Sifre, Simon Schmitt, Arthur Guez, Edward Lockhart, Demis Hassabis, Thore Graepel, Timothy Lillicrap, and David Silver. Mastering atari, go, chess and shogi by planning with a learned model. *Nature*, 588: 604–609, 2020.
- David Silver, Julian Schrittwieser, Karen Simonyan, Ioannis Antonoglou, Aja Huang, Arthur Guez, Thomas Hubert, Lucas Baker, Matthew Lai, Adrian Bolton, Yutian Chen, Timothy Lillicrap, Fan Hui, Laurent Sifre, George van den Driessche, Thore Graepel, and Demis Hassabis. Mastering the game of go without human knowledge. *Nature*, 550:354–359, 2017.
- David Silver, Thomas Hubert, Julian Schrittwieser, Ioannis Antonoglou, Matthew Lai, Arthur Guez, Marc Lanctot, Laurent Sifre, Dhharshan Kumaran, Thore Graepel, Timothy Lillicrap, Karen Simonyan, and Demis Hassabis. A general reinforcement learning algorithm that masters chess, shogi, and go through self-play. *Science*, 362(6419):1140–1144, 2018.
- Brian D. Ziebart, Andrew Maas, J. Andrew Bagnell, and Anind K. Dey. Maximum entropy inverse reinforcement learning. In *AAAI Conference on Artificial Intelligence*, 2008.